

Multithreading

Josh Wilcox (jw14g24)

May 11, 2025

Table of Contents

① Creating and Using Threads in Java

- Extends Thread Class
- Implements Runnable Interface
- Thread Properties and Methods

② States and Life Cycles of Threads

- States of a Java Thread
- Useful Methods for Java Threads
- Life Cycle of a thread

③ Synchronisation

- Synchronising in Java
 - Example of using Multiple Locks
- Deadlocks

④ Communicating Between threads

Creating and Using Threads in Java

Why use threads

- Efficiency through concurrency
 - Facilitates using multiple core CPUs that are available
- Makes computer programs simpler to right
 - Separating complex tasks into parallel operations
 - Handling multiple user interactions simultaneously
- Prevents the inefficiency caused by sequential processes when they can be parallelisable

Threading in the JVM

- In the JVM, each thread has its own **Private Memory Space** to store its own Java Stack and a program counter
- The metaspace (where method details are stored) and the heap are **shared between threads**

The Methods of Creating Threads

- There are two main ways to create threads in Java:
 - `class MyThread extends Thread`
 - `class MyRunner implements Runnable`

Extends Thread Class

Information on the thread Class

- The `Thread` class is defined in `java.lang`
- Any class that extends the `Thread` class can be created as its own thread
 - This is done by **overriding** the `run()` method which details how the main process the thread should achieve
 - **IMPORTANT** - `run()` does not actually start the class! To start it use `start()`
 - * `run()` is automatically called by `start()`
 - * `start()` only allows the thread to start if it has the right status

Example

```
1  public class MyThread extends Thread {  
2      private String name;  
3      public MyThread(String name) {  
4          this.name = name;  
5      }  
6  
7      public void run() {  
8          for (int i=0; i<5; i++) {  
9              System.out.println(name + " is executing, i = " + i);  
10         }  
11     }  
12 }  
13  
14 public class ThreadDemo01 {  
15     public static void main (String[] args) {  
16         MyThread mt1 = new MyThread("Thread A");  
17         MyThread mt2 = new MyThread("Thread B");  
18         // Start each thread - They will now run in parallel  
19         mt1.start();  
20         mt2.start();  
21     }  
22 }
```

Implements Runnable Interface

Information on the Runnable Interface

- The **Runnable** interface only has one abstract method:
 - `public abstract void run();`
- This means that a new class that **implements** the **Runnable** class does **not** have its own **start()** class manually
- So we still do use the **Thread** class in order to start the thread.

```
1 public class MyThread implements Runnable {  
2     @Override  
3     public void run(){  
4         // Run Logic Here  
5     }  
6     public static void main(String[] args){  
7         // Create the class with the run logic  
8         MyThread myThread = new MyThread();  
9         // Encapsulate the class in a Thread Object  
10        Thread threadRunner = new Thread(myThread);  
11        // Start the Thread  
12        threadRunner.start();  
13    }  
14}
```

Why use a Runnable Interface over a Thread Abstract Class

- A class can implement **Multiple** interfaces
- A class can only extend **One** abstract class
- Therefore using the **Runnable** interface allows our threads to both extend an abstract class and implement other interfaces
 - This provides more opportunity for polymorphism

Thread Methods

Thread.sleep(long n)

- Makes the current thread pause for n milliseconds
- During this time, the thread enters a "sleeping" state
- Must be in a try-catch block as it throws InterruptedException

```
1 try {
2     Thread.sleep(1000); // Sleep for 1 second
3 } catch (InterruptedException e) {
4     e.printStackTrace();
5 }
```

Thread Priorities

- Threads have priorities from 1 (lowest) to 10 (highest)
- Set using `setPriority(int priority)`
- Constants available: MIN_PRIORITY, NORM_PRIORITY, MAX_PRIORITY
- Higher priority threads are generally executed first

Thread.yield()

- Temporarily pauses current thread to let other threads execute
- Thread moves from "running" state to "ready" state
- No guarantee which thread will execute next
- Primarily used for thread scheduling optimization
 - Helps prevent any single thread from monopolizing CPU time
 - Particularly useful in busy-wait scenarios
 - Can improve overall system responsiveness
- Important considerations:
 - Behavior is highly platform-dependent
 - Only a scheduling hint - may be ignored by the JVM
 - Not recommended for synchronization purposes
- Example usage:

```
1 if(threadNeedsToWait) {
2     Thread.yield(); // Give other threads a chance
3 }
```

Interrupting Threads

- Use `interrupt()` to interrupt a thread
- Interrupted threads throw InterruptedException
- Check if interrupted using `isInterrupted()`
- Useful for stopping long-running or sleeping threads

States and Names of a Java Thread

What are the States of a Thread

- A thread in Java exist in any of six states
- Each state is stored as a key in an `Enum` object
 - NEW - Thread that has not yet started
 - RUNNABLE - Thread that is **executing** in the JVM
 - BLOCKED - Thread that is vlocked and waiting for a **monitor lock**
 - WAITING - Thread is waiting **indefinitely** for another thread to perform a particular action
 - TIMED_WAITING - Like WAITING, but includes a maximum waiting time
 - TERMINATED - The Thread has exited

Setters and Getters for Thread Names

- We can use constructors to define names for our threads
 - `public Thread (String name)`
 - `public Thread(Runnable target, String name)`
- Normally set threads names before execution, but we can change them afterwards using `setName(String name)`
- You can get the name using `getName()`

Daemon Threads

- Deamon threads are low-priority threads whose only role is to **provide services to user threads**
 - User threads are the normal threads we are used to
- Main Properties of Daemon Threads:
 - Won't prevent JVM from exiting while running
 - * Normally, JVM only terminates when **all user threads** have finished executing
 - Threads created in a daemon threads will also be daemon
- Use `setDaemon(boolean on)` to true to make a thread a daemon

Useful Methods for Java Threads

isAlive()

- Determines whether a thread has been started or has been terminated
- Returns true if the thread has been started and has not yet died
- Returns false if either:
 - The thread hasn't been started yet
 - The thread has completed its execution
 - The thread has been terminated

join()

- The `join()` method is used to make one thread wait for another thread's completion
- When `threadA.join()` is called from `threadB`:
 - `threadB` will pause its execution
 - `threadB` will wait until `threadA` completes
 - This ensures sequential execution when needed
- Has three variants:
 - `join()` - waits indefinitely
 - `join(long millis)` - waits for specified milliseconds
 - `join(long millis, int nanos)` - adds nanosecond precision
- Common use cases:
 - Waiting for background tasks to complete
 - Ensuring proper order of operations
 - Coordinating dependent thread activities

Life Cycle of a thread

Thread States and Transitions

- New State
 - Thread is created but not yet started
 - Transitions to Runnable when `start()` is called
- Runnable State
 - Thread is executing or ready to execute
 - Can transition to:
 - * Running - when selected by scheduler
 - * Blocked/Waiting - when requesting resources
 - * Terminated - when execution completes
- Blocked/Waiting States
 - Thread is temporarily inactive
 - Caused by:
 - * `sleep()` calls
 - * I/O operations
 - * Lock acquisition
 - * `join()` calls
- Terminated State
 - Thread has completed execution
 - Cannot be restarted

Synchronisation

The need for Synchronisation

- All threads in the JVM **share the heap**
 - Any fields in an object can be accessible to a number of threads at the same time
- After a thread has updated a field with a value, it's possible another thread comes along and updates the same field with a different value
- When one thread is trying to read the value from a field, another thread may be trying to change that value
- This can lead to a lot of funkiness to do with value change order
- Even simple operations like incrementing are not atomic:
 - A statement like `z = z + 1` compiles to multiple bytecode instructions
 - These include loading from heap, incrementing on stack, and storing back
 - Threads can be interrupted between any of these steps
 - This creates race conditions even for seemingly atomic operations

Incrementation Example

- See below, the final output for `z` should be 4

```
1 public class UnsafeCounter {
2     private int z = 2; // Initial value is 2
3     void incr() {
4         int x = z; // Thread 1: reads z=2
5             // Thread 2: reads z=2 (before Thread 1 updates z)
6         x = x + 1; // Thread 1: x becomes 3
7             // Thread 2: x becomes 3 (both threads working with original z value)
8         z = x; // Thread 1: sets z=3
9             // Thread 2: sets z=3 (overwrites Thread 1's update)
10    } // Final z=3, even though we wanted z=4
11 }
12 public class ThreadDemo {
13     public static void main(String[] args) {
14         UnsafeCounter counter = new UnsafeCounter();
15         Thread t1 = new Thread(() -> counter.incr());
16         Thread t2 = new Thread(() -> counter.incr());
17         t1.start();
18         t2.start();
19         // After both threads finish, z might still be 3 instead of 4!
20     }
21 }
```

Synchronising and Locks in Java

Locks

- If some thread is holding the lock of a synchronised block or method, no other thread can access that block!
- We can use locks to ensure **only one thread** can access code at any one point
 - This kind of lock is known as **mutex lock** which provides **mutual exclusion**
- Every object in Java has a mutex lock associated with it

Synchronised Blocks

- You can **lock** a defined section of code using a **synchronised block**:

```
1 synchronized(lockObject){  
2     // Code Here...  
3 }
```

Synchronised Methods

- Methods can be marked as synchronised using the **synchronized** keyword
- This automatically synchronises the entire method body
- Example:

```
1 public synchronized void increment() {  
2     count++; // This operation is now thread-safe  
3 }
```

- Key points about synchronised methods:
 - Only one thread can execute a synchronized method at a time
 - The lock is on the entire object for instance methods
 - For static methods, the lock is on the Class object

The problem with Synchronisation

- Using synchronised blocks is very inefficient in many circumstances
- There are heavy runtime overheads in using locks and blocking threads
 - Try to synchronise the **minimum** number of instructions you can

'this' lock and multiple locks

- Using `synchronized(this)` or declaring a method as `synchronized` locks the entire object, preventing any other synchronized method/block on the same object from executing concurrently.
- For finer control, use dedicated lock objects (e.g., `private final Object lock = new Object();`) to synchronize only specific data or operations.
- Static synchronized methods lock on the `Class` object, so they do not interfere with instance-level locks.
- Multiple locks can be used within a class to protect independent resources, increasing concurrency and reducing contention.
- Always document which locks protect which data to avoid confusion and bugs.

Synchronising in Java - Example of using Multiple Locks

Synced Bank account using Sync Blocks

```
1  public class BankAccount {
2      private int savingsBalance = 0;
3      private int checkingBalance = 0;
4
5      // Separate lock objects for each account type
6      private final Object savingsLock = new Object();
7      private final Object checkingLock = new Object();
8
9      // Deposit to savings account
10     public void depositToSavings(int amount) {
11         synchronized (savingsLock) {
12             savingsBalance += amount;
13         }
14     }
15
16     // Deposit to checking account
17     public void depositToChecking(int amount) {
18         synchronized (checkingLock) {
19             checkingBalance += amount;
20         }
21     }
22
23     // Get savings balance
24     public int getSavingsBalance() {
25         synchronized (savingsLock) {
26             return savingsBalance;
27         }
28     }
29
30     // Get checking balance
31     public int getCheckingBalance() {
32         synchronized (checkingLock) {
33             return checkingBalance;
34         }
35     }
36 }
```

- This example allows threads to safely update or read the savings and checking balances independently.
- A thread updating the savings account does not block another thread updating the checking account, increasing concurrency.

Deadlocks

What is a Deadlock?

- A deadlock occurs when two or more threads are waiting for each other to release resources
- Each thread holds a resource that another thread needs
- None can proceed because they're all waiting
- Real-life example:
 - Two cars meet on a narrow bridge from opposite directions
 - Neither can proceed because both need the other to back up
 - Both are waiting for the other to move first

Code Example of Deadlock

```
1 public class DeadlockExample {  
2     private Object lock1 = new Object();  
3     private Object lock2 = new Object();  
4  
5     public void method1() {  
6         synchronized(lock1) {  
7             System.out.println("Method 1 has lock1");  
8             try { Thread.sleep(100); } catch(InterruptedException e) {}  
9             synchronized(lock2) {  
10                 System.out.println("Method 1 has lock2");  
11             }  
12         }  
13     }  
14  
15     public void method2() {  
16         synchronized(lock2) {  
17             System.out.println("Method 2 has lock2");  
18             try { Thread.sleep(100); } catch(InterruptedException e) {}  
19             synchronized(lock1) {  
20                 System.out.println("Method 2 has lock1");  
21             }  
22         }  
23     }  
24 }
```

Preventing Deadlocks

- Always acquire locks in the same order
- Use timeouts when acquiring locks
- Avoid nested locks where possible
- Use `tryLock()` methods from `java.util.concurrent.locks`
- Implement deadlock detection mechanisms

Communicating Between threads

The need for communication handling

- Sometimes threads need to communicate or coordinate
 - For example, a consumer thread may need to **wait** for a producer to put data in the shared buffer
 - Java has methods that allow threads to pause execution until a certain criterion is met
 - * Threads can also **notify** other threads when such condition has changed

A note on Object Monitors

- A monitor is a mechanism that ensures:
 - Only one thread can execute a synchronized method/block at a time
 - Other threads must wait for the active thread to finish
 - Provides a way for threads to coordinate through wait/notify
- When a thread enters a synchronized block:
 - It acquires the object's monitor
 - Other threads cannot acquire the same monitor
 - Monitor is released when thread exits the synchronized block

wait

- The **wait()** method causes **the thread it is called in** to release the lock it holds and **wait** until another thread calls **notify()** or **notifyAll()**
- Must be called inside a **synchronized block or method**
- Can also include a timeout by adding numbers in milliseconds to the constructor
- The **wait()** method is called on the same object as the lock in the synchronised block

notify

- The **notify()** method wakes up a **single** thread that has previously been **waiting** on the object's **monitor**
- The awakened thread only proceed if the synchronised block is **unlocked** - it will still wait until this is true before proceeding
- The **notify()** method is called on the same object as the lock in the synchronised block

notifyAll

- The **notifyAll()** method has the same core functionality as **notify()** but wakes up **all** threads waiting on the object's **monitor**
- Only one thread will acquire the lock and proceed; others will continue waiting for the lock.
- The **notifyAll()** method is called on the same object as the lock in the synchronised block